

SECTION 0 — INTERVIEW STRATEGY (Senior SDET Masterclass)

Topics Covered

- 0.1 Deloitte Role Expectations (Senior Consultant / Specialist Senior)
- 0.2 Senior SDET Competency Matrix (Junior vs Mid vs Senior vs Lead)
- 0.3 Interview-Round Strategy (The 5-Round Deloitte Loop)
- 0.4 Technical vs Scenario-Based Questions (Frameworks & Spoken Scripts)
- 0.5 Project Deep-Dive Strategy (The C-A-S-C-I Master Narrative & Architecture)
- 0.6 Coding-Round Strategy (The C-E-B-O-T Narration System)
- 0.7 Automation Architecture Strategy (Layered Decoupling & Boundary Rules)
- 0.8 Behavioral & Leadership Questions (Consulting & Client Ownership)
- 0.9 STARI Answer Structure (Situation-Task-Action-Result-Insight)
- 0.10 Strategic Questions to Ask the Interviewer (Panel, Manager, Partner)
- 0.11 Final Interview Checklist (24h, 1h, During, and Artifact Portfolio)
- 0.12 Current Baseline & 14-Day Remediation Roadmap

0.1 Deloitte Role Expectations

Deloitte hires Senior SDETs primarily into **Senior Consultant** or **Specialist Senior** tracks (under Core Business Operations, Quality Engineering, or Cloud Engineering practices). At this level, you are evaluated as a **billable, client-facing engineering authority and technical owner**, not a script executor who waits for test cases.

DELOITTE SENIOR SDET CORE ROLE PILLARS		
1. ARCHITECTURE Zero-leak drivers, isolated fixtures, scalable runners	2. SHIFT-LEFT CI/CD quality gates, ephemeral test environments	3. CONSULTING ADVISORY Business risk translation, scope governance, pushback on impossible deadlines
4. CAPABILITY BUILDING & MENTORING Code reviews, PR quality standards, onboarding accelerators, CoE assets		

The 4 Pillars of Senior Consulting Delivery

1. **End-to-End Automation Architecture:** You must design and own complete automation ecosystems (Playwright/TypeScript or Selenium/Java/TestNG, REST Assured, and SQL/DB integration) from scratch. You own separation of concerns, thread-safe concurrency, dynamic data generation, resilient locator contracts, and distributed execution.
2. **Shift-Left CI/CD Governance:** You own the quality gates within GitHub Actions, Jenkins, or Azure DevOps. You understand how to turn a 2-hour blocking regression suite into an 8-minute sharded PR gate, enforce flaky

test quarantine budgets, and manage containerized test runners (Docker/Testcontainers).

3. **Consulting Advisory & Client Influence:** You operate inside client pods with ambiguity. You translate technical defects into business risk (\$ lost, SLA breach penalties, customer churn). When clients ask for "100% test automation in 2 weeks," you don't say yes; you present a risk-weighted automation ROI matrix.
4. **Talent Multiplier & Mentoring:** You establish PR standards, author framework style guides, conduct code reviews, and upskill junior engineers and manual testers into automation contributors.

Key Metrics Interviewers Listen For

In your answers, you must anchor your achievements to industry-standard engineering and delivery metrics:

- **Flaky Test Ratio (SLO):** Reduced from 15–20% to `< 2%` across a 3,000+ test suite.
- **Pipeline Feedback Loop:** PR feedback time reduced from 45–60 minutes down to `< 10 minutes`.
- **Defect Detection Percentage (DDP):** In-sprint automation catching `> 85%` of defects before staging.
- **Escaped Defect Rate (EDR):** Production escapes reduced by `> 60%` year-over-year.
- **DORA Metrics:** Deployment frequency accelerated from bi-weekly to daily; Change Failure Rate (CFR) cut to `< 5%`; Mean Time to Recovery (MTTR) dropped to `< 45 minutes`.

Trap Points & Anti-Patterns to Avoid

- **✗ "I automate user stories assigned to me by the QA Lead."** (Signals junior order-taker. Instead say: "I analyzed the sprint backlog, mapped user journeys to risk tiers, and designed the automation suite for critical paths.")
- **✗ "Our goal was 100% automation coverage with zero defects."** (Signals lack of enterprise reality. Instead say: "We targeted an 80/20 risk-based coverage model, prioritizing critical revenue paths on UI and pushing comprehensive edge cases to API and contract layers.")
- **✗ "The developers wrote buggy code and refused to write unit tests."** (Deloitte tests for collaborative client empathy. Instead say: "We had a testability gap, so I partnered with the dev lead to introduce contract testing and agreed upon stable `data-testid` attributes as part of the Definition of Done.")

0.2 Senior SDET Competency Matrix

Competency Dimension	Junior SDET (1–3 yrs)	Mid-Level SDET (3–5 yrs)	Senior SDET (5–8+ yrs) — Deloitte Target	Lead / Principal SDET (8+ yrs)
Language & DSA	Basic syntax, loops, uses ArrayList/HashMap without depth.	Understands OOP, basic Collections, writes linear solutions.	Low-level JVM memory (Stack vs Heap, GC tuning), thread safety, Big-O tradeoffs, stream pipelines.	Language internals, bytecode, custom classloaders, high-throughput memory optimizations.
UI Automation	Records scripts, uses <code>Thread.sleep()</code> , brittle XPath.	Writes POM, uses explicit waits, basic TestNG runner.	Web-first assertions, auto-waiting mechanics, CDP/DevTools protocols, shadow DOM, iframe boundaries.	Framework design from scratch, self-healing architecture, cross-browser engine benchmarking.

Competency Dimension	Junior SDET (1–3 yrs)	Mid-Level SDET (3–5 yrs)	Senior SDET (5–8+ yrs) — Deloitte Target	Lead / Principal SDET (8+ yrs)
API Automation	Executes Postman collections, basic REST Assured GET/POST.	Asserts status codes and basic JSON fields using JsonPath.	Schema validation, OAuth2 token caching/refresh, request/response specs, POJO Jackson builders, WireMock.	Enterprise contract testing (Pact), distributed tracing (OTel), gRPC/GraphQL automation, chaos testing.
Concurrency & Scale	Runs tests sequentially on local machine.	Configures <code>thread-count</code> in <code>testng.xml</code> but suffers flaky state.	<code>ThreadLocal</code> lifecycle management, worker isolation in Playwright, database race condition prevention.	Dynamic grid autoscaling on K8s (Selenoid/KEDA), distributed sharding across 50+ cloud runners.
CI/CD & DevOps	Triggers manual builds in Jenkins UI.	Edits basic Jenkinsfile or GitHub Actions YAML steps.	Designs multi-stage declarative pipelines, matrix sharding, artifact caching, failure reporting, PR gates.	GitOps pipeline architecture, Docker image optimization, ephemeral preview environment provisioning.
Data Strategy	Hardcodes test data in scripts or Excel sheets.	Uses Faker libraries or static JSON files.	API-based pre-seeding + teardown, worker-scoped data isolation, database transaction rollback.	Synthetic data generation engines, GDPR/HIPAA data masking pipelines, event-driven state hydration.
Triage & Debugging	Re-runs failed tests manually until they pass.	Inspects stack traces and screenshots.	Systematic triage: inspects HAR/network, browser console logs, DOM snapshots, time-travel traces.	Flakiness analytics dashboards, automated quarantine pipelines, root-cause clustering algorithms.

0.3 Interview-Round Strategy (The 5-Round Deloitte Loop)

DELOITTE SENIOR SDET INTERVIEW LOOP

[Round 1] Recruiter / Talent Acquisition (15-20 min)

└ Filter: Role alignment, communication, notice, stack match

[Round 2] Technical Assessment / Screening (60 min)

└ Filter: Live coding (Java/TS), DSA, SQL queries, MCQs

[Round 3] Core Technical & Automation Deep Dive (60 min)

└ Filter: Framework architecture, waits, API, multithreading
[Round 4] Systems, Scenarios & Whiteboarding (60 min)
└ Filter: Triage broken CI, scale to 5k tests, design from scratch
[Round 5] Partner / Director Techno-Behavioral (45 min)
└ Filter: Client leadership, consulting mindset, culture fit

Answering Strategy: The 60 / 30 / 10 Delivery Rule

Whenever answering a technical or architectural question, divide your verbal response into three deliberate phases:

1. **60% — Engineering Solution:** Direct, technically precise explanation using exact terminology (protocols, memory states, concurrency models).
2. **30% — Trade-offs & Constraints:** Why you chose this approach over alternatives (e.g., *Playwright fixtures vs beforeEach*, *API seeding vs UI setup*, *ThreadLocal vs synchronization*).
3. **10% — Business & Scale Outcome:** Concrete metric impact (*"This reduced suite runtime by 68% and eliminated 99% of data collisions across 16 parallel workers"*).

0.4 Technical vs Scenario-Based Questions

1. Technical Questions (Testing Depth & Mechanics)

- **Goal:** Evaluates whether you truly understand the engine or merely memorized syntax.
- **Delivery Pattern:**
 1. *One-Sentence Architectural Definition.*
 2. *Under-the-Hood Mechanics* (network packets, memory allocation, DOM lifecycle).
 3. *Enterprise Scale Caveat / Gotcha.*

Verbal Example: "Why do you prefer Web-First Assertions over generic assertions?"


"Generic assertions like `Jest expect(await page.isVisible())` take a static boolean snapshot at one millisecond in time; if the element is mid-transition or re-rendering via React, it immediately fails without retrying. Web-first assertions like Playwright's `await expect(locator).toBeVisible()` invert this control: they continuously poll the DOM, re-evaluating the locator and actionability criteria until the expectation passes or the timeout expires (default 5s). This single architectural shift eliminates arbitrary `sleep()` statements and removes ~80% of timing flakiness in single-page applications."

2. Scenario Questions (Testing Judgment & Triage)

- **Goal:** Evaluates your composure and systematic troubleshooting when production or releases are on the line.
- **The SIRH Framework:**
 - **S — Signal & Quarantine:** How do you contain the failure immediately so the team isn't blocked?

- **I — Isolate & Reproduce:** How do you separate test code, application code, test data, and infrastructure?
- **R — Remediate:** What is the architectural fix (not just a band-aid)?
- **H — Harden & Prevent:** What automated guardrail (linter, rule, gate) guarantees this never recurs?

Scenario Triage Playbook: "Your test passes locally on Mac but fails 20% of the time in headless Linux CI. Walk me through your triage."



Verbal Script for Interview:

"I follow a 4-step triage methodology:

1. First, I don't guess—I retrieve the Playwright `trace.zip` or Selenium video/logs recorded on failure from CI artifacts. In Trace Viewer, I inspect the exact action, DOM snapshot, console errors, and network waterfall at the point of failure.
2. I check for Environment Disparities: Local machines are fast and headed; CI runners (e.g. GitHub Actions Linux VMs) have restricted vCPUs (2 cores), slower disk I/O, and default to 1280x720 headless viewports, which frequently causes responsive menus to collapse into hamburger buttons. I ensure CI runs at 1920x1080 and pins OS font packages.
3. I verify Concurrency & Data Isolation: If tests run with 4 workers in CI, tests that passed serially locally might be colliding on shared database records or user logins. I ensure every test derives its data from `workerIndex` or UUIDs.
4. Once root cause is identified (e.g., element hidden behind an animation), I fix it with an auto-retrying web-first locator and add a lint rule banning `waitForTimeout` or hard sleeps across the repo."

0.5 Project Deep-Dive Strategy (The C-A-S-C-I Master Narrative)

In Round 3 and Round 4, you will be asked: "Walk me through the test automation framework you built in your most recent project."

Use the **C-A-S-C-I** narrative (Context $\rightarrow$ Architecture $\rightarrow$ Scale $\rightarrow$ Challenge $\rightarrow$ Impact).

1. Spoken 2-Minute Elevator Pitch

"In my recent engagement with a major digital banking and financial services platform, I served as the Senior Automation Architect owning the test engineering workstream across 4 agile pods.

When I joined, the team was burdened by a 4-hour manual regression suite and an inherited legacy Selenium Java framework that suffered from a 22% flakiness rate in CI, causing developers to completely ignore pipeline red builds.

*I designed and delivered a greenfield, hybrid automation ecosystem using **Playwright with TypeScript** for modern web channels, coupled with **REST Assured / Axios API clients** for fast state hydration, and containerized **PostgreSQL** verification hooks. Architecturally, we enforced strict layer separation: raw locators and browser interactions were strictly encapsulated inside Page and Component objects; business actions returned immutable representations; and test classes only expressed user intent and assertions.*


To scale this, we integrated our suite into **GitHub Actions with 8-way matrix sharding, pre-authenticated test contexts using `storageState`, and API-based test data factories that generated dynamic users per worker.*


The biggest architectural hurdle was handling third-party payment gateway callbacks and WebSocket notifications without introducing hard waits. We solved this by implementing network route interception and mock fallbacks for external dependencies while polling internal Kafka events using event listener fixtures.

As a result, we compressed our total regression cycle from **4 hours down to 11 minutes, dropped our flaky test rate from **22% to under 1.5%**, and accelerated release velocity from monthly drops to twice-weekly on-demand production deployments with zero P1 defect escapes across three quarters.**

2. Framework Architectural Whiteboard Diagram

When asked to whiteboard your architecture, sketch this exact multi-tiered diagram:





EXECUTION & OBSERVABILITY INFRA

- | | |
|---|----------------------------------|
| • Matrix Sharded CI (GitHub Actions / Jenkins Pipeline) | • Trace Viewer (`trace.zip`) |
| • Docker Container Runners | • Slack / Teams Failure Webhooks |
| • Allure / Playwright HTML Reports | |

0.6 Coding-Round Strategy (The C-E-B-O-T System)

During the live coding round, interviewers evaluate your **problem-solving hygiene, edge-case vigilance, and verbal reasoning**, not just working code.

C-E-B-O-T CODING PROTOCOL

- | | |
|------------------|---|
| [C] CLARIFY | Ask input constraints, null policy, scale (2 min) |
| [E] EDGE CASES | List empty, single, duplicates, overflow (1 min) |
| [B] BRUTE FORCE | State naive $O(N^2)$ approach aloud before coding (1 min) |
| [O] OPTIMIZE | Identify bottleneck; state optimal $O(N)$ data structure |
| [T] TEST DRY-RUN | Trace sample input with pointers before saying "done" |

Verbal Script for Narrating Big-O & Trade-Offs

"To solve this problem of finding the first non-repeating character in a string:

- The **brute force approach** is to run nested loops comparing every character against all others, which yields $O(N^2)$ time complexity and $O(1)$ space complexity.
- To **optimize this for enterprise scale** (e.g., parsing large server log streams), we can trade space for time. We can use a single pass with an integer array of size 26 or a `LinkedHashMap` to store frequencies and preserve insertion order.
- This optimizes time complexity to $O(N)$ with $O(1)$ auxiliary space (since the alphabet size is bounded to 26/128 characters). Let me write this out with null and empty string guardrails first."

0.7 Automation Architecture Strategy

The 4 Non-Negotiable Architectural Rules

1. The Downward-Only Dependency Rule:

- Tests depend on Pages/Components .

- Pages/Components depend on Core Framework Utilities .
- Core Framework Utilities depend on the underlying driver (Playwright / WebDriver).
- **Crucial:** Infrastructure layers NEVER depend on test cases. Tests NEVER bypass Pages to execute raw driver commands.

2. Zero Assertions in Page Objects:

- Page Objects represent the *state and actions* of the system under test; they must return data, locators, or new Page Objects.
- Assertions belong exclusively inside the test methods (`spec.ts` or `@Test`). Putting assertions inside Page Objects prevents reusability across positive, negative, and exploratory test flows.

3. Stateless Test Isolation:

- Tests must be completely hermetic. Test A must never rely on data generated by Test B.
- Every test must establish its own preconditions via fast API endpoints or isolated fixtures and clean up its state on teardown.

4. Thread-Safe Driver & Context Lifecycle:

- **In Selenium/Java:** You must encapsulate WebDriver inside a `ThreadLocal<WebDriver>` and explicitly call `tlDriver.remove()` inside `@AfterMethod` to avoid thread pool memory leaks in long-running CI runners.
- **In Playwright:** You must leverage built-in worker process boundaries and ephemeral `BrowserContext` fixtures instead of creating global browser singletons.

0.8 Behavioral & Leadership Questions (Consulting Context)

Deloitte scores candidates on **Client Leadership, Delivery Ownership, and Integrity**.

1. Disagreement with Developer on Defect Severity

- **Question:** "A developer rejects your P1 defect, claiming it is an edge case that users won't hit. How do you handle it?"
- **What Evaluator Scores:** Objective data over emotion; customer and business risk framing.
- **Spoken Script:**

"I never turn defect triage into a subjective debate. First, I verify reproducibility by attaching complete evidence: network logs, payload payloads, browser console logs, and step-by-step video. Next, I review production telemetry and user analytics (e.g., Datadog or Google Analytics) to show the actual percentage of users utilizing that specific browser or checkout flow. If the business impact represents revenue loss or compliance breach, I schedule a 10-minute sync with the Developer and Product Owner, framing the decision around risk acceptance: 'If we ship with this defect, here is the quantified exposure (\$X revenue / Y users). Does the business accept this risk, or should we patch it?' The Product Owner makes the business call, and our engineering partnership remains collaborative and respectful."

2. Pushing Back on Unrealistic Client Deadlines

- **Question:** "The client demands 100% automation of 500 complex test cases in 2 weeks. How do you respond?"
- **What Evaluator Scores:** Consulting maturity, scope negotiation, risk management.

- **Spoken Script:**

"I don't simply say 'no'; I present an evidence-based roadmap. I explain that rushing 500 automated tests in two weeks leads to brittle, unmaintainable scripts that fail constantly, destroying confidence in the automation suite. Instead, I conduct a Risk-Based Coverage Analysis:

- We identify the top 20% of critical revenue and security paths (the smoke suite of ~40-50 tests) and automate those immediately with enterprise-grade stability.
- We leverage API automation for the underlying business logic, which delivers 5x faster coverage than UI automation.
- We build a transparent 6-week burn-down chart showing prioritized, phased automation deliverables. The client gets reliable, high-ROI quality gates immediately without inheriting technical debt."

0.9 STARI Answer Structure

For behavioral questions, structure every response using **STARI** (Situation $\rightarrow$ Task $\rightarrow$ Action $\rightarrow$ Result $\rightarrow$ Insight) and deliver it in under **90 seconds**:

0s	15s	30s	70s	85s	90s
Situation	Task	Action	Result	Insight	
Context,	Your	3 technical & leadership	Numbers,	Lasting	
scale &	specific	steps (what you did, how	metrics,	impact	
stakes	ownership	you led, why it worked)	%, \$, time	& rule	

Complete STARI Model: The Flaky Suite Recovery Story

- **Situation (15s):** "At my previous engagement, our 1,200-test regression suite had deteriorated to an 18% flaky failure rate in CI. The engineering team had lost complete trust in the pipeline, and developers routinely clicked 'Re-run' 3 to 4 times until builds passed."
- **Task (15s):** "As the Senior SDET Lead, I took full ownership of recovering pipeline integrity, reducing flakiness below 2%, and establishing a strict quality gate within 30 days."
- **Action (40s):**
 1. "First, I implemented an automated **Flaky Quarantine Pipeline**: any test failing intermittently across consecutive runs was tagged `@quarantine` and routed to a non-blocking diagnostic job with full video and trace capture, keeping the main PR build green."
 2. "Second, I clustered the root causes: 60% were due to shared test data collisions across parallel workers, 30% were race conditions on dynamic SPA elements, and 10% were third-party API timeouts. I replaced all shared database seeds with UUID-based API factories and refactored brittle assertions to Playwright web-first auto-retrying matchers."
 3. "Third, I instituted a team-wide 'Flakiness Budget' SLO: no squad could merge new features if their module's flake rate exceeded 2%."
- **Result (15s):** "Within four weeks, our flaky test rate dropped from **18% to 1.1%**. CI build runtimes decreased from **55 minutes to 14 minutes**, developer re-runs fell to zero, and the team saved an estimated 120 engineering hours per month."

- **Insight (5s):** *"I learned that test flakiness is primarily an architectural and governance problem, not a timing problem. Enforcing strict data isolation and quarantine budgets is the only way to sustain long-term pipeline trust."*
-

0.10 Strategic Questions to Ask the Interviewer

Always close the interview with strategic, high-value questions that position you as an architectural leader:

For Technical Panelists / SDET Leads

1. *"How do you currently handle test data management and environment teardown across parallel CI runs—do you rely on ephemeral containerized databases, API seeding, or dedicated staging databases?"*
2. *"What does your flakiness triage workflow look like today? Do you have an automated quarantine pipeline with an SLA, or are flaky tests addressed ad-hoc during release cycles?"*
3. *"What is the current distribution in your test pyramid between unit, API contract, and end-to-end UI tests, and where is the biggest strategic gap you are looking to close?"*

For Engineering Managers / Delivery Leads

4. *"What is the expectation around Definition of Done and testability—do developers actively collaborate on adding stable test locators (`data-testid`) and writing integration tests, or does the SDET workstream operate after code freeze?"*
5. *"What does success look like for this Senior SDET role in the first 90 days? What measurable outcomes or delivery milestones will prove that this hire was a success?"*

For Partners / Practice Directors (Deloitte Leadership)

6. *"How is Deloitte's Quality Engineering practice incorporating generative AI and agentic workflows into client test frameworks, and what opportunities exist to contribute to reusable CoE assets?"*
 7. *"On large enterprise transformation projects, how does your leadership team manage client expectations when legacy systems lack testability or API contracts?"*
-

0.11 Final Interview Checklist

24 Hours Before Interview

- ☐ Rehearse the **2-minute C-A-S-C-I project pitch** out loud until delivery is effortless.
- ☐ Memorize your 5 core impact metrics: **Suite size (3,500)**, **Runtime (4h \$to\$ 11m)**, **Flake rate (18% \$to\$ 1.2%)**, **PR feedback (<10m)**, **Defect escapes (0 P1s)**.
- ☐ Practice drawing the **4-tier Framework Architecture diagram** on a blank sheet of paper in under 3 minutes.
- ☐ Warm up on live coding: String manipulation, HashMap frequency counting, and two-pointer arrays.
- ☐ Review your resume line-by-line: Be prepared to defend every tool, library, and configuration mentioned.

1 Hour Before Interview

- ☐ Test screen share, audio, video, and IDE settings (VS Code / IntelliJ with dark theme and legible font size).
- ☐ Open a blank scratchpad for taking notes on interviewer problem statements.
- ☐ Keep water nearby and review your 3 STARI behavioral stories.

During the Interview

- ☐ Use the **60 / 30 / 10 rule** on technical questions.
- ☐ Never jump straight into code: apply **C-E-B-O-T** (Clarify $\rightarrow$ Edge Cases $\rightarrow$ Brute Force $\rightarrow$ Optimize $\rightarrow$ Test).
- ☐ On framework design, prioritize **state isolation, thread safety, and maintenance costs** over syntax tricks.

0.12 Current Baseline & 14-Day Remediation Roadmap

Rate your mastery honestly from 1 to 5 (1 = Aware, 3 = Project Proficient, 5 = Can Architect & Teach):

14-DAY SENIOR SDET SPRINT PLAN	
Days 1 – 3	Core Java, Memory Layouts, Collections & Concurrency
Days 4 – 6	Playwright & TypeScript Architecture Deep Dive
Days 7 – 8	REST Assured, API Client Design & Contract Testing
Days 9 – 10	CI/CD Pipelines, Matrix Sharding & Docker Runners
Days 11 – 12	Complex SQL Queries, Window Functions & DB Testing
Days 13 – 14	Full Mock Loops: Architecture Whiteboarding & STARI

Technical Domain	Self Score (1–5)	Senior Interview Bar	Immediate Action Item
Java Foundations & OOP	[]	Explain JVM Stack/Heap, ClassLoaders, GC algorithms, SOLID in automation.	Review Section 1, 2, and 3.
Collections & HashMap	[]	Bitwise bucket index $(n-1) \& \text{hash}$, treeification, Java 8 ConcurrentHashMap CAS.	Review Section 4 and 5.
Concurrency & Threading	[]	ThreadPoolExecutor tuning, ThreadLocal lifecycle, race condition diagnosis.	Review Section 9.
Playwright & TypeScript	[]	Auto-waiting mechanics, custom fixtures, network route mocking, storageState.	Review Section 19.
API & REST Assured	[]	Request/Response specifications, POJO builders, OAuth2 caching, WireMock stubs.	Review Section 18 and 25.
CI/CD & DevOps	[]	Multi-stage YAML, matrix sharding, artifact caching, flaky quarantine gates.	Review Section 21.

Technical Domain	Self Score (1–5)	Senior Interview Bar	Immediate Action Item
SQL & Data Persistence	[]	Complex joins, window functions (<code>DENSE_RANK</code>), transaction rollbacks.	Review Section 24.
Framework Architecture	[]	Clean separation of concerns, zero assertions in POM, downward-only dependencies.	Review Section 20 and 31.
Project Story & STARI	[]	Flawless 2-minute elevator pitch with 5 quantifiable metrics.	Practice Section 0 and 34.