

## SECTION 19 — PLAYWRIGHT WITH TYPESCRIPT (Senior SDET Masterclass)

### Topics Covered


- **19.1 Playwright Architecture & Mechanics** (DevTools/CDP/WebSocket pipe vs Selenium HTTP REST)
- **19.2 Browser, BrowserContext, and Page Model** (Process boundaries, incognito contexts in <50ms)
- **19.3 Configuration & Projects Matrix** ( `playwright.config.ts` , mobile emulation, multi-browser)
- **19.4 Test Isolation & Fixtures Architecture** ( `test.extend` , worker fixtures, auto fixtures)
- **19.5 Resilient Locators** ( `page.getByRole` , `page.getByTestId` , `page.getByLabel` , strict mode)
- **19.6 Auto-Waiting Engine & The 6 Actionability Checks** (Attached, Visible, Stable, Enabled, Editable, Receives Events)
- **19.7 Web-First Assertions** ( `expect(locator).toBeVisible()` polling vs generic sync assertions)
- **19.8 Multi-Tab, Popups, and iframe ( `frameLocator` ) Automation**
- **19.9 StorageState & Authentication Session Reuse** (Login once, share cookies/tokens across workers)
- **19.10 Network Interception & API Mocking** ( `page.route` , `route.fulfill` , `route.continue` , modifying payloads)
- **19.11 APIRequestContext** (Direct Node.js HTTP client without browser for data seeding)
- **19.12 Parallelism & Sharding** ( `fullyParallel` , `--workers` , `--shard=1/4` , merging HTML reports)
- **19.13 Observability: Time-Travel Tracing** ( `trace.zip` , Trace Viewer), Screenshots, and Videos
- **19.14 Component Object Model in TypeScript** (Production Architecture)
- **19.15 High-Stakes Senior Playwright Interview Questions & Spoken Solutions**

### 19.1 Playwright Architecture & Mechanics

#### Architectural Paradigm: Playwright vs. Selenium WebDriver

Selenium WebDriver relies on a multi-hop, stateless HTTP REST architecture defined by the W3C WebDriver specification. Every browser interaction is serialized into an individual HTTP request, routed through a standalone out-of-process driver daemon ( `chromedriver` , `geckodriver` ), and translated into browser commands.

Playwright abandons the stateless HTTP model entirely. It establishes a **single, persistent, bidirectional multiplexed WebSocket / OS pipe connection** between the Node.js test runner client and an out-of-process Playwright driver process.



[Chrome Browser Engine]

## PLAYWRIGHT ARCHITECTURE

[Playwright Test Runner (Node.js Process)]

JSON-RPC 2.0 Multiplexed Protocol over Single Pipe / WebSocket  
 Zero HTTP handshake overhead; Bidirectional event streaming

[Playwright Driver (Node.js C++ Binding Core)]

▼ Chrome DevTools (CDP)  
 [Chromium / Blink Engine]

▼ Juggler Protocol  
 [Firefox / Gecko Engine]

▼ WebKit Protocol  
 [WebKit Engine]

## Protocol Internals & Transport Layers

### 1. Multiplexed Bidirectional Connection:

- Rather than opening a new TCP connection or sending individual HTTP POST requests per action, Playwright communicates via an internal **JSON-RPC protocol over standard input/output pipes (IPC)** or a local **WebSocket**.
- Both browser commands ( `click` , `navigate` ) and browser-emitted telemetry ( `console` , `pageerror` , `network requests` , `DOM mutations` , `frame transitions` ) share the same persistent channel.

### 2. Engine Instrumentation:

- For **Chromium**, Playwright connects directly to the Chrome DevTools Protocol (CDP) session ( `Target.attachToTarget` , `DOMWorld` , `Runtime.evaluate` ).
- For **Firefox** and **WebKit**, Playwright does not use WebDriver or CDP. Microsoft maintains upstream-instrumented browser binaries:
  - **Firefox**: Instrumented with **Juggler**, an internal debugging protocol providing native event hooks directly in SpiderMonkey/Gecko.
  - **WebKit**: Instrumented with WebKit's native remote inspector protocol, patched to support frame switching, network interception, and pointer emulation identical to Blink.

### 3. Execution Latency & Throughput:


- Selenium HTTP round-trip latency: ~15ms to 50ms per command due to HTTP request parsing, TCP headers, and driver response serialization.
- Playwright JSON-RPC pipelining: <1ms per command over local IPC stream. Batching and event streaming allow Playwright to listen for DOM mutations reactively rather than polling via polling loops.

## 19.2 Browser, BrowserContext, and Page Model

Playwright structures browser hierarchy into three distinct boundaries: `Browser` , `BrowserContext` , and `Page` .

BROWSER INSTANCE (Heavy OS Process, e.g., chrome.exe)

- Launches browser engine, allocates shared GPU/render processes, manages sockets



## Memory Layout and Context Lifecycle

- **Browser** :
  - Represents the physical OS process (e.g., Chromium with browser, GPU, and utility processes).
  - Heavyweight: Consumes 150MB–300MB RAM. Spawning takes 1,000ms–2,500ms.
  - In Playwright Test, a single **Browser** instance is launched **once per worker process** and reused across hundreds of tests.
- **BrowserContext** :
  - An ephemeral, completely isolated browser session analogous to a distinct Incognito window.
  - **Zero Process Overhead**: Creating a **BrowserContext** does **not** fork a new operating system process. It registers an isolated in-memory profile partition inside the browser engine.
  - **Instantiation Time**: Created and destroyed in **10ms to 30ms**.
  - **Complete Isolation**: Cookies, **localStorage**, **sessionStorage**, IndexedDB, service workers, and HTTP cache are strictly partitioned.
- **Page** :
  - A single tab or window within a **BrowserContext**. Multiple pages in the same context share cookies and authentication tokens.

## Multi-Context Enterprise Testing Pattern

Simulating multi-user interactions (e.g., real-time chat, buyer-seller marketplaces, administrative privilege escalation) without opening two browser windows:

```
import { test, expect } from '@playwright/test';

test('Real-time collaborative chat between buyer and seller', async ({ browser }) => {
  // Create isolated session for Buyer
  const buyerContext = await browser.newContext({
    storageState: '.auth/buyer.json',
    viewport: { width: 1280, height: 720 },
  });
```

```

const buyerPage = await buyerContext.newPage();

// Create isolated session for Seller
const sellerContext = await browser.newContext({
  storageState: '.auth/seller.json',
  viewport: { width: 1280, height: 720 },
});
const sellerPage = await sellerContext.newPage();

// Buyer sends an offer
await buyerPage.goto('/listing/item-492');
await buyerPage.getByRole('button', { name: /make offer/i }).click();
await buyerPage.getByLabel(/offer amount/i).fill('450');
await buyerPage.getByRole('button', { name: /submit offer/i }).click();

// Seller receives real-time notification via WebSocket
await sellerPage.goto('/dashboard/inbox');
const offerItem = sellerPage.getByRole('listitem').filter({ hasText: 'Item 492' });
await expect(offerItem).toContainText('$450.00');
await offerItem.getByRole('button', { name: /accept/i }).click();

// Buyer verifies offer acceptance confirmation
await expect(buyerPage.getByRole('status')).toContainText('Offer Accepted');

// Teardown contexts cleanly
await buyerContext.close();
await sellerContext.close();
});

```

## 19.3 Configuration & Projects Matrix ( playwright.config.ts )

A production `playwright.config.ts` manages environment variables, dependency chains, cross-browser matrices, mobile viewports, network settings, and CI sharding parameters.

```

import { defineConfig, devices } from '@playwright/test';
import * as path from 'path';

export const STORAGE_STATE_PATH = path.join(__dirname, '.auth/user.json');

export default defineConfig({
  testDir: './tests',
  fullyParallel: true,
  forbidOnly: !!process.env.CI,
  retries: process.env.CI ? 2 : 0,
  workers: process.env.CI ? '100%' : undefined,

  // Enterprise Timeout Architecture
  globalTimeout: 60 * 60 * 1000, // 1 hour maximum pipeline runtime
  timeout: 45 * 1000, // 45 seconds per individual test execution
  expect: {
    timeout: 7 * 1000, // 7 seconds web-first assertion polling
  },
});

```

```

reporter: process.env.CI
  ? [
    ['blob', { outputDir: 'blob-report' }],
    ['junit', { outputFile: 'results/junit.xml' }],
    ['github'],
  ]
  : [['html', { open: 'on-failure' }]],

use: {
  baseURL: process.env.BASE_URL || 'https://staging.enterprise.internal',
  trace: 'on-first-retry',
  screenshot: 'only-on-failure',
  video: 'retain-on-failure',
  actionTimeout: 10 * 1000,
  navigationTimeout: 15 * 1000,
  ignoreHTTPSErrors: true,
},

projects: [
  // Global Authentication Setup Project
  {
    name: 'setup-auth',
    testMatch: /\.*\setup\.ts/,
  },

  // Desktop Browsers (Dependent on setup-auth)
  {
    name: 'chromium-desktop',
    use: {
      ...devices['Desktop Chrome'],
      storageState: STORAGE_STATE_PATH,
    },
    dependencies: ['setup-auth'],
  },
  {
    name: 'firefox-desktop',
    use: {
      ...devices['Desktop Firefox'],
      storageState: STORAGE_STATE_PATH,
    },
    dependencies: ['setup-auth'],
  },
  {
    name: 'webkit-desktop',
    use: {
      ...devices['Desktop Safari'],
      storageState: STORAGE_STATE_PATH,
    },
    dependencies: ['setup-auth'],
  },

  // Mobile Emulation Projects
  {
    name: 'mobile-chrome',
    use: {
      ...devices['Pixel 7'],

```

```

        storageState: STORAGE_STATE_PATH,
      },
      dependencies: ['setup-auth'],
    },
    {
      name: 'mobile-safari',
      use: {
        ...devices['iPhone 14 Pro'],
        storageState: STORAGE_STATE_PATH,
      },
      dependencies: ['setup-auth'],
    },
  ],

  // Automatic Local Server Lifecycle Management
  webServer: process.env.CI ? undefined : {
    command: 'npm run start:mock-server',
    port: 3000,
    timeout: 120 * 1000,
    reuseExistingServer: true,
  },
});

```

## 19.4 Test Isolation & Fixtures Architecture

### Dependency Injection vs. Shared Mutable State

Legacy automation frameworks (JUnit, TestNG) rely on `@BeforeMethod` and mutable class variables ( `protected WebDriver driver;` ). In concurrent execution, thread crosstalk, incomplete cleanup, and static references cause widespread test flakiness.

Playwright features a **hierarchical Dependency Injection (DI) system** via `test.extend<T>()` . Fixtures are:

- **On-Demand:** A fixture runs only if a test explicitly includes it in its parameter signature.
- **Composable:** Fixtures can depend on other fixtures.
- **Hermetic:** Tear-down logic executes automatically inside the `use()` statement boundary even if test steps throw exceptions.

#### FIXTURE EXECUTION TIMELINE

|                     |                                                 |
|---------------------|-------------------------------------------------|
| [1. Setup Phase]    | Allocates resource (e.g., API seed, auth state) |
| [2. Injection]      | <code>await use(initializedResource);</code>    |
| [3. Test Execution] | Test runs and performs actions / assertions     |
| [4. Teardown Phase] | Execution returns to fixture; runs cleanup      |

### Production Implementation: Custom Fixture with API Seeding & StorageState

```
// fixtures/enterprise-fixtures.ts
import { test as base, expect, Page, APIRequestContext } from '@playwright/test';

export interface UserSessionData {
  userId: string;
  email: string;
  token: string;
  organizationId: string;
}

type EnterpriseFixtures = {
  authenticatedPage: Page;
  testUserData: UserSessionData;
};

type EnterpriseWorkerFixtures = {
  adminApiClient: APIRequestContext;
};

export const test = base.extend<EnterpriseFixtures, EnterpriseWorkerFixtures>({
  // Worker-scoped fixture: Initialized once per OS worker process
  adminApiClient: [async ({ playwright }, use) => {
    const apiContext = await playwright.request.newContext({
      baseURL: process.env.API_GATEWAY_URL || 'https://api.staging.enterprise.internal',
      extraHTTPHeaders: {
        'Authorization': `Bearer ${process.env.ADMIN_SERVICE_ROLE_KEY}`,
        'Content-Type': 'application/json',
      },
    },
  ]);

  await use(apiContext);

  // Teardown worker context
  await apiContext.dispose();
}, { scope: 'worker' }]);

// Test-scoped fixture: Seeds isolated dynamic user per test
testUserData: async ({ adminApiClient }, use) => {
  const uniqueSuffix = Date.now() + Math.random().toString(36).substring(2, 7);
  const testUser = {
    email: `sdet_test_${uniqueSuffix}@enterprise.internal`,
    password: 'SecureEnterprisePass123!',
    organizationName: `Org_${uniqueSuffix}`,
  };

  // Fast-path API creation
  const res = await adminApiClient.post('/v1/admin/users', { data: testUser });
  if (!res.ok()) {
    throw new Error(`Failed to seed test user: ${res.status()} ${await res.text()}`);
  }
  const userPayload = await res.json();

  const sessionData: UserSessionData = {
    userId: userPayload.id,
    email: testUser.email,
    token: userPayload.sessionToken,
  };
}
```

```

    organizationId: userPayload.orgId,
  };

  // Inject seeded data into test
  await use(sessionData);

  // Hermetic Teardown: Guaranteed cleanup after test completes
  await adminApiClient.delete(`/v1/admin/users/${sessionData.userId}`);
},

// Test-scoped fixture: Provides Page pre-loaded with authenticated state
authenticatedPage: async ({ browser, testUserData }, use) => {
  const context = await browser.newContext({
    baseURL: 'https://staging.enterprise.internal',
  });

  // Seed session token into localStorage before page loads
  await context.addInitScript((token) => {
    window.localStorage.setItem('AUTH_TOKEN', token);
  }, testUserData.token);

  const page = await context.newPage();
  await use(page);

  await context.close();
},
});

export { expect };

```

## 19.5 Resilient Locators & Strict Mode

### The Locator Hierarchy

Playwright discourages implementation-dependent CSS and XPath selectors ( `div > span:nth-child(3)` ). Instead, it prioritizes **accessible, user-facing semantics** that mirror how real users and screen readers perceive the interface.

| Locator Strategy                 | Method                                                    | Accessibility / Robustness Rationale                                                                         |
|----------------------------------|-----------------------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| <b>ARIA Role (Gold Standard)</b> | <code>page.getByRole('button', { name: 'Submit' })</code> | Aligns test automation with accessibility trees; resilient to styling refactors.                             |
| <b>Form Label</b>                | <code>page.getByLabel('Work Email')</code>                | Locates input elements associated via <code>&lt;label for="..."&gt;</code> or <code>aria-labelledby</code> . |
| <b>Text Content</b>              | <code>page.getByText('Invoice #1042')</code>              | Direct text matching; supports substring, regex, and exact matching.                                         |
| <b>Placeholder</b>               | <code>page.getByPlaceholder('Search products...')</code>  | Locates inputs lacking visible labels.                                                                       |

| Locator Strategy | Method                                                  | Accessibility / Robustness Rationale                                                        |
|------------------|---------------------------------------------------------|---------------------------------------------------------------------------------------------|
| Explicit Test ID | <code>page.getByTestId('checkout-order-summary')</code> | Stable engineering contract ( <code>data-testid</code> ), isolated from UI/styling changes. |

Strict Mode Mechanics

By default, every Playwright locator operates in **Strict Mode**. If a selector expression matches **more than one** DOM element, any subsequent action ( `click()` , `fill()` , `hover()` ) immediately halts execution and throws an informative `StrictnessError` .

```
// Throws Error if multiple 'Submit' buttons exist in the DOM
await page.getByRole('button', { name: 'Submit' }).click();
// Error: locator.click: Error: strict mode violation: getByRole('button', { name: 'Submit' }) r
// 1) <button class="btn-primary">Submit</button> aka getByRole('button', { name: 'Submit' }).
// 2) <button class="btn-secondary">Submit</button> aka getByRole('button', { name: 'Submit' })
```


Locator Chaining & Filtering (Composition over Fragile Paths)

```
// Scope down by component boundaries using locator.filter
const invoiceRow = page.getByRole('row').filter({
  has: page.getByRole('cell', { name: 'INV-2026-9901' }),
  hasNot: page.getByRole('status', { name: 'Paid' }),
});

// Perform action within strictly filtered scope
await invoiceRow.getByRole('button', { name: /process payment/i }).click();
```

19.6 Auto-Waiting Engine & The 6 Actionability Checks

Playwright eliminates the need for manual synchronization ( `Thread.sleep()` , arbitrary polling delays). Before executing any action, Playwright's core engine evaluates the **Actionability Checks Matrix**.



directly to element (not intercepted by spinner/modal)

### Actionability Check Matrix by Action Type

| Action                                          | Attached | Visible | Stable | Enabled | Editable | Receives Events |
|-------------------------------------------------|----------|---------|--------|---------|----------|-----------------|
| <code>click()</code>                            | YES      | YES     | YES    | YES     | NO       | YES             |
| <code>fill()</code> / <code>clear()</code>      | YES      | YES     | NO     | YES     | YES      | NO              |
| <code>check()</code> / <code>unchecked()</code> | YES      | YES     | YES    | YES     | NO       | YES             |
| <code>hover()</code>                            | YES      | YES     | YES    | NO      | NO       | YES             |
| <code>press()</code> / <code>type()</code>      | YES      | YES     | NO     | YES     | NO       | NO              |
| <code>scrollIntoViewIfNeeded()</code>           | YES      | NO      | NO     | NO      | NO       | NO              |

[!CAUTION]

#### The Pitfall of `{ force: true }`:

Passing `{ force: true }` to `.click()` disables checks 2 through 6, dispatching a synthetic JavaScript synthetic event directly to the element. If an overlay modal or loading spinner covered the button, real users could not click it, but your test would pass—introducing critical false-positive test results.

## 19.7 Web-First Assertions

### Synchronous vs. Web-First Retrying Assertions

In traditional frameworks, assertions evaluate state at a single instantaneous point in time:

```
// ANTI-PATTERN: Instantaneous snapshot assertion (Inherently Flaky in SPAs)
const isVisible = await page.getByRole('alert').isVisible();
expect(isVisible).toBe(true); // Fails if alert renders 50ms after check!
```

Playwright introduces **Web-First Assertions** (`await expect(locator)...`). These assertions do not take a point-in-time snapshot. They continuously re-poll the browser engine until the condition matches or the assertion timeout (default: 5,000ms) expires.

```
// PRODUCTION STANDARD: Continuous Web-First Assertion
await expect(page.getByRole('alert')).toBeVisible({ timeout: 10_000 });
```

### Essential Web-First Assertion Methods

```
// 1. Text & Content Matching (Auto-retrying)
await expect(page.getByRole('heading', { level: 1 })).toHaveText('Enterprise Dashboard');
await expect(page.getByTestId('status-pill')).toContainText(/active/i);

// 2. Attributes and DOM Properties
await expect(page.getByRole('button', { name: 'Save' })).toBeEnabled();
await expect(page.getByRole('checkbox', { name: 'Agree' })).toBeChecked();
await expect(page.getByRole('textbox', { name: 'Email' })).toHaveValue('sdet@company.com');
await expect(page.getByRole('link', { name: 'Docs' })).toHaveAttribute('target', '_blank');

// 3. Count Verification (Guards against partial renders)
await expect(page.getByRole('listitem')).toHaveCount(10);

// 4. Inverted Negative Assertions (Auto-waits for element disappearance)
await expect(page.getByRole('progressbar')).not.toBeVisible();

// 5. Soft Assertions (Continues test execution to collect full failure report)
await expect.soft(page.getByTestId('user-first-name')).toHaveText('Jane');
await expect.soft(page.getByTestId('user-last-name')).toHaveText('Doe');
```

## 19.8 Multi-Tab, Popups, and iframe ( `frameLocator` ) Automation

### Multi-Tab & Popup Handling via Event Synchronization

When clicking a link with `target="_blank"` or an OAuth button, race conditions often cause tests to fail if the child page opens before the test attaches a listener. Playwright uses a **concurrent `Promise.all` idiom**:

```
test('Handle OAuth Login Popup Window', async ({ page }) => {
  await page.goto('/login');

  // Concurrently register event listener and trigger popup click
  const [popupPage] = await Promise.all([
    page.waitForEvent('popup'),
    page.getByRole('button', { name: /sign in with google/i }).click(),
  ]);

  // popupPage is a fully functional Page instance within the same BrowserContext
  await popupPage.waitForLoadState('domcontentloaded');
  await popupPage.getByLabel(/email or phone/i).fill('enterprise_user@gmail.com');
  await popupPage.getByRole('button', { name: /next/i }).click();

  // Once popup window closes, parent page automatically reflects auth state
  await popupPage.close();
  await expect(page.getByRole('heading', { name: /welcome back/i })).toBeVisible();
});
```

### iframe Encapsulation with `frameLocator`

Unlike Selenium's imperative `driver.switchTo().frame()`, Playwright provides declarative, auto-waiting `frameLocator` encapsulation that seamlessly handles nested iframes.

```

test('Process Embedded Stripe Payment inside Nested iframes', async ({ page }) => {
  await page.goto('/checkout');

  // Declarative frame locator chain: Parent frame -> Child card input frame
  const paymentContainer = page.frameLocator('iframe[name="stripe-checkout"]');
  const cardElementFrame = paymentContainer.frameLocator('iframe[title="Secure card number input

  // Actions inside the frame apply the full 6 Actionability Checks automatically
  await cardElementFrame.getByPlaceholder('Card number').fill('4242424242424242');

  const expiryFrame = paymentContainer.frameLocator('iframe[title="Secure expiration date input
  await expiryFrame.getByPlaceholder('MM / YY').fill('12/28');

  // Submit order button resides in top-level DOM
  await page.getByRole('button', { name: /complete order/i }).click();
  await expect(page.getByRole('alert')).toHaveText('Payment Authorized');
});

```

## 19.9 StorageState & Authentication Session Reuse

### The Authentication Bottleneck at Scale

Running 5,000 end-to-end tests where every test logs in through the UI (filling credentials, solving CAPTCHA/MFA, processing redirects) introduces massive overhead:

$$\text{Overhead} = 5,000 \text{ tests} \times 4.5 \text{ seconds/login} = 22,500 \text{ seconds} \approx 6.25 \text{ hours}$$

Playwright resolves this via **StorageState Caching**: execute authentication **once** in a dedicated setup project, dump the session state (cookies, `localStorage`) to a JSON file, and hydrate every parallel worker context instantly (<10ms).

### Implementation: Authentication Setup Project

```

// tests/auth.setup.ts
import { test as setup, expect } from '@playwright/test';
import * as path from 'path';

const authFile = path.join(__dirname, '../.auth/admin.json');

setup('Authenticate Admin and Persist Storage State', async ({ page }) => {
  // Option A: UI-based single login
  await page.goto('/login');
  await page.getByLabel(/username/i).fill(process.env.ADMIN_USERNAME!);
  await page.getByLabel(/password/i).fill(process.env.ADMIN_PASSWORD!);
  await page.getByRole('button', { name: /sign in/i }).click();

  // Wait until session is established
  await expect(page.getByRole('navigation')).toBeVisible();

  // Capture all cookies, localStorage tokens, and origins to disk

```

```
await page.context().storageState({ path: authFile });
});
```

## Bypassing UI with API-Hydrated StorageState

For ultimate efficiency, bypass the UI entirely and hydrate `storageState` directly via API inside the setup phase:

```
setup('Authenticate directly via REST API and write storageState', async ({ playwright, request } => {
  const response = await request.post('https://api.staging.enterprise.internal/v1/auth/login', {
    data: {
      username: process.env.ADMIN_USERNAME,
      password: process.env.ADMIN_PASSWORD,
    },
  });
  expect(response.ok()).toBeTruthy();
  const { accessToken, refreshToken, user } = await response.json();

  // Create empty context, populate storage state, and write to disk
  const context = await playwright.request.newContext();
  const storageState = {
    cookies: [
      {
        name: 'session_token',
        value: refreshToken,
        domain: '.enterprise.internal',
        path: '/',
        expires: -1,
        httpOnly: true,
        secure: true,
        sameSite: 'Lax' as const,
      },
    ],
    origins: [
      {
        origin: 'https://staging.enterprise.internal',
        localStorage: [
          { name: 'ACCESS_TOKEN', value: accessToken },
          { name: 'USER_PROFILE', value: JSON.stringify(user) },
        ],
      },
    ],
  };

  const fs = await import('fs/promises');
  await fs.writeFile(authFile, JSON.stringify(storageState, null, 2));
});
```

## 19.10 Network Interception & API Mocking

Playwright allows granular interception, aborting, and mutating of HTTP/HTTPS requests directly at the browser network layer before packets hit the wire.

## NETWORK ROUTING PIPELINE

[Browser Page Request]



[page.route('\*\*/api/v1/invoices', route =&gt; { ... })]

- ├─ route.abort('failed') → Drop packet instantly
- ├─ route.fulfill({ json }) → Return mock payload
- └─ route.continue() / fetch() → Intercept & mutate real API

## Production Examples: Mocking, Chaos Engineering, and Asset Blocking

```
test('Simulate 500 Internal Server Error & UI Error Boundary', async ({ page }) => {
  // Mock API response with 500 status code
  await page.route('**/api/v1/billing/summary', async (route) => {
    await route.fulfill({
      status: 500,
      contentType: 'application/json',
      body: JSON.stringify({
        code: 'BILLING_SERVICE_UNAVAILABLE',
        message: 'Database connection pool exhausted',
      }),
    });
  });

  await page.goto('/dashboard/billing');

  // Verify UI displays graceful fallback error boundary
  await expect(page.getByRole('alert')).toContainText('Unable to retrieve billing summary');
  await expect(page.getByRole('button', { name: /retry/i })).toBeVisible();
});

test('Mutate Live API Response (Upstream Response Interception)', async ({ page }) => {
  // Fetch real upstream response, mutate payload, and fulfill
  await page.route('**/api/v1/user/entitlements', async (route) => {
    // 1. Fetch real response from actual backend
    const response = await route.fetch();
    const json = await response.json();

    // 2. Mutate real data to inject VIP beta flag
    json.features.push('AI_BETA_PREVIEW');
    json.tier = 'ENTERPRISE_PLUS';

    // 3. Fulfill route with modified payload
    await route.fulfill({
      response,
      json,
    });
  });

  await page.goto('/settings/features');
  await expect(page.getByTestId('ai-preview-banner')).toBeVisible();
});
```

```
});

test('Enterprise Performance: Block Third-Party Analytics Trackers', async ({ page }) => {
  // Drastically speeds up test execution and isolates test from network noise
  await page.route(
    /(analytics\.google\.com|datadoghq-browser-agent|sentry\.io|hotjar\.com)/,
    (route) => route.abort()
  );

  await page.goto('/checkout');
  // Tests execute cleanly without third-party analytics telemetry delays
});
```

## 19.11 APIRequestContext

Playwright provides `APIRequestContext` —a direct, headless HTTP client that runs entirely within the Node.js test process, without spawning or orchestrating a browser instance.

### Architecture of Hybrid Fast-Path Testing

- **Setup:** Seed database records, create test organizations, and assign roles via `APIRequestContext` in **100ms**.
- **Execution:** Navigate directly to target URL and validate critical user paths in the browser.
- **Teardown:** Issue direct `DELETE` requests via `APIRequestContext` to clean up resources immediately.

```
import { test, expect } from '@playwright/test';

test('Fast-path hybrid test: API seed + UI verify', async ({ page, request }) => {
  // 1. Seed customer record via headless API client (< 80ms)
  const createCustomerRes = await request.post('/api/v1/customers', {
    data: {
      name: 'Acme Global Logistics',
      creditLimit: 750000,
      tier: 'PLATINUM',
    },
    headers: {
      'Authorization': `Bearer ${process.env.INTERNAL_SERVICE_TOKEN}`,
    },
  });
  expect(createCustomerRes.ok()).toBeTruthy();
  const customer = await createCustomerRes.json();

  // 2. Launch UI directly into the created customer's view
  await page.goto(`/customers/${customer.id}/overview`);
  await expect(page.getByRole('heading', { level: 1 })).toHaveText('Acme Global Logistics');
  await expect(page.getByTestId('credit-limit')).toHaveText('$750,000.00');

  // 3. Fast-path teardown
  const deleteRes = await request.delete(`/api/v1/customers/${customer.id}`);
  expect(deleteRes.status()).toBe(204);
});
```

## 19.12 Parallelism & Sharding

### Worker Model & Process Isolation

Playwright test runners execute tests across independent **Node.js worker processes**:

- Each worker process runs its own isolated Node runtime.
- Workers do not share memory, variables, or state.
- `fullyParallel: true` enables tests within the same file to run concurrently across separate workers.
- If a test fails, Playwright can discard the worker process and spin up a clean worker to prevent state leakage during retries.

#### PLAYWRIGHT TEST SCHEDULER

|                                  |             |                     |
|----------------------------------|-------------|---------------------|
| Worker Process 1 (Node PID 4010) | → Browser 1 | → [Test A] [Test C] |
| Worker Process 2 (Node PID 4011) | → Browser 2 | → [Test B] [Test D] |
| Worker Process 3 (Node PID 4012) | → Browser 3 | → [Test E] [Test G] |
| Worker Process 4 (Node PID 4013) | → Browser 4 | → [Test F] [Test H] |

### CI Sharding across Distributed Runners

For suites with thousands of tests, split execution horizontally across ephemeral CI nodes using `--shard=x/y`.

```
# Matrix execution across 4 parallel CI runners
# Machine 1:
npx playwright test --shard=1/4 --reporter=blob
# Machine 2:
npx playwright test --shard=2/4 --reporter=blob
# Machine 3:
npx playwright test --shard=3/4 --reporter=blob
# Machine 4:
npx playwright test --shard=4/4 --reporter=blob
```

### GitHub Actions Matrix & Report Merge Pipeline

```
name: Enterprise Playwright E2E Suite
on:
  push:
    branches: [main]
  pull_request:
    branches: [main]

jobs:
  test-sharded:
    timeout-minutes: 45
    runs-on: ubuntu-latest
    strategy:
      fail-fast: false
      matrix:
```

```

    shardIndex: [1, 2, 3, 4]
    shardTotal: [4]
  steps:
    - uses: actions/checkout@v4
    - uses: actions/setup-node@v4
      with:
        node-version: 20
        cache: 'npm'
    - run: npm ci
    - run: npx playwright install --with-deps

    - name: Run Playwright Shard
      run: npx playwright test --shard=${{ matrix.shardIndex }}/${{ matrix.shardTotal }}
      env:
        CI: true

    - name: Upload Shard Blob Report
      if: always()
      uses: actions/upload-artifact@v4
      with:
        name: blob-reports-${{ matrix.shardIndex }}
        path: blob-report/
        retention-days: 1

merge-reports:
  if: always()
  needs: [test-sharded]
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v4
    - uses: actions/setup-node@v4
      with:
        node-version: 20
    - run: npm ci

    - name: Download all Shard Reports
      uses: actions/download-artifact@v4
      with:
        path: all-blob-reports
        pattern: blob-reports-*
        merge-multiple: true

    - name: Merge HTML Reports
      run: npx playwright merge-reports --reporter=html ./all-blob-reports

    - name: Publish Consolidated HTML Report
      uses: actions/upload-artifact@v4
      with:
        name: final-playwright-report
        path: playwright-report/
        retention-days: 14

```

## 19.13 Observability: Time-Travel Tracing, Screenshots, Videos

## Anatomy of `trace.zip`

Playwright Tracing produces a self-contained ZIP archive containing rich execution telemetry:

- `trace.trace` : Monotonically timestamped JSON-RPC protocol events, internal actions, and errors.
- `trace.network` : Full HTTP archive (HAR) recording headers, bodies, cookies, and timings for every network transaction.
- `trace.stacks` : Precise TypeScript call stacks mapped via source maps to exact lines of code.
- Snapshots: Periodic DOM snapshots (serialized before and after every action) and screencast video frames.

```

trace.zip
├── 0.png, 1.png, 2.png      # Screencast timeline thumbnails
├── trace.trace              # Action log and actionability timing checks
├── trace.network           # Comprehensive network request/response HAR
└── resources/              # CSS stylesheets, DOM snapshots, HTML frames

```

## Trace Viewer Inspection

View captured traces locally or in the browser:

```
npx playwright show-trace path/to/trace.zip
```

Inside Trace Viewer, SDETs can:

1. Scrub through the visual timeline before and after individual actions.
2. Inspect the live DOM state at the exact moment of a click.
3. Verify actionability check logs to see which check stalled (e.g., waiting for element to receive pointer events).
4. Inspect request payloads, console messages, and source code.

## Enterprise Observability Configuration

```

// playwright.config.ts
export default defineConfig({
  use: {
    // Collect traces only on first retry to avoid storage bloat on passing runs
    trace: 'on-first-retry',
    // Capture screenshot only on failure
    screenshot: 'only-on-failure',
    // Retain video recording only on failure
    video: 'retain-on-failure',
  },
});


```

## 19.14 Component Object Model (COM) in TypeScript

### Page Objects vs. Component Objects

Traditional Page Object Models (POM) often grow into bloated god-classes containing hundreds of locators across an entire page.

The **Component Object Model (COM)** decomposes pages into modular, composable, reusable UI components. A component accepts a parent `Locator` and scopes all internal elements relative to that parent, making it reusable across multiple pages.



## Production Implementation: Component-Driven Architecture

### 1. Base Component Class

```
// components/base.component.ts
import { Locator, Page } from '@playwright/test';

export abstract class BaseComponent {
  protected readonly root: Locator;
  protected readonly page: Page;

  constructor(root: Locator) {
    this.root = root;
    this.page = root.page();
  }

  async isVisible(): Promise<boolean> {
    return this.root.isVisible();
  }
}
```

### 2. Reusable Data Table Component

```
// components/data-table.component.ts
import { Locator, expect } from '@playwright/test';
import { BaseComponent } from '../base.component';

export class TableRowComponent extends BaseComponent {
  readonly cells: Locator = this.root.getByRole('cell');
```

```

    async getCellText(index: number): Promise<string> {
        return (await this.cells.nth(index).innerText()).trim();
    }

    async clickAction(buttonName: string | RegExp): Promise<void> {
        await this.root.getByRole('button', { name: buttonName }).click();
    }
}

export class DataTableComponent extends BaseComponent {
    readonly headerRow: Locator = this.root.getByRole('row').first();
    readonly rows: Locator = this.root.getByRole('row').filter({
        hasNot: this.root.page().getByRole('columnheader'),
    });

    getRow(index: number): TableRowComponent {
        return new TableRowComponent(this.rows.nth(index));
    }

    getRowByText(text: string | RegExp): TableRowComponent {
        const matchedRow = this.rows.filter({ hasText: text });
        return new TableRowComponent(matchedRow);
    }

    async expectRowCount(expected: number): Promise<void> {
        await expect(this.rows).toHaveCount(expected);
    }
}

```

### 3. Page Object Composing Components

```

// pages/invoice-management.page.ts
import { Page, Locator, expect } from '@playwright/test';
import { DataTableComponent } from '../components/data-table.component';

export class InvoiceManagementPage {
    readonly page: Page;
    readonly table: DataTableComponent;
    readonly searchInput: Locator;
    readonly statusFilterSelect: Locator;

    constructor(page: Page) {
        this.page = page;
        this.searchInput = page.getByPlaceholder('Search invoices...');
        this.statusFilterSelect = page.getByRole('combobox', { name: /filter by status/i });

        // Instantiate component scoped to the data table container
        this.table = new DataTableComponent(page.getByRole('table', { name: /invoices table/i }));
    }

    async navigate(): Promise<void> {
        await this.page.goto('/invoices');
        await expect(this.page.getByRole('heading', { level: 1, name: 'Invoices' })).toBeVisible();
    }
}

```

```

    async filterByStatus(status: 'Pending' | 'Paid' | 'Overdue'): Promise<void> {
      await this.statusFilterSelect.selectOption(status);
    }

    async voidInvoice(invoiceId: string): Promise<void> {
      const row = this.table.getRowByText(invoiceId);
      await row.clickAction(/void invoice/i);

      // Confirm in modal dialog
      const modal = this.page.getByRole('dialog', { name: /confirm void/i });
      await expect(modal).toBeVisible();
      await modal.getByRole('button', { name: /confirm/i }).click();
      await expect(modal).not.toBeVisible();
    }
  }
}

```

#### 4. Type-Safe Test Execution

```

// tests/invoices.spec.ts
import { test, expect } from '../fixtures/enterprise-fixtures';
import { InvoiceManagementPage } from '../pages/invoice-management.page';

test.describe('Enterprise Invoice Workflow', () => {
  test('Search and void pending invoice using Component Objects', async ({ authenticatedPage }) {
    const invoicePage = new InvoiceManagementPage(authenticatedPage);

    await invoicePage.navigate();
    await invoicePage.filterByStatus('Pending');
    await invoicePage.searchInput.fill('INV-2026-4402');

    await invoicePage.table.expectRowCount(1);
    const row = invoicePage.table.getRow(0);
    expect(await row.getCellText(1)).toBe('INV-2026-4402');

    await invoicePage.voidInvoice('INV-2026-4402');
    await expect(authenticatedPage.getByRole('status')).toContainText('Invoice voided successful');
  });
});

```

## 19.15 High-Stakes Senior Playwright Interview Questions & Spoken Solutions

**Question 1: How does Playwright's architecture differ under the hood from Selenium WebDriver, and why does Playwright execute significantly faster with less flakiness?**

**Spoken Answer:**

"Selenium relies on a multi-hop, stateless HTTP REST architecture conforming to the W3C WebDriver specification. Every test command ( `findElement` , `click` ) is serialized into an individual HTTP request, sent

over a TCP socket to an out-of-process driver daemon like `chromedriver`, which translates it and sends it to the browser. This creates considerable round-trip latency—roughly 15 to 40 milliseconds per command—and because HTTP is stateless, the driver cannot reactively listen to browser DOM mutations without explicit polling loops.

In contrast, Playwright establishes a single, multiplexed, bidirectional WebSocket or OS pipe connection directly to the browser process using JSON-RPC. For Chromium, it talks directly to the Chrome DevTools Protocol; for Firefox and WebKit, Microsoft maintains instrumented engine builds with internal protocols like Juggler.

This architecture provides three major advantages:

1. **Sub-millisecond latency:** Commands stream over a persistent pipe with zero HTTP handshake overhead.
2. **Bidirectional event streaming:** Playwright listens directly to browser lifecycle events—DOM mutations, network requests, frame navigations, and web workers—allowing it to react immediately rather than polling.
3. **The Auto-Waiting Engine:** Playwright executes 6 built-in actionability checks (Attached, Visible, Stable, Enabled, Editable, Receives Events) before every action. It verifies that CSS animations have finished and hit-testing resolves to the element before dispatching clicks. This fundamentally eliminates the timing race conditions that cause flakiness in Selenium."

**Question 2: In a suite of 5,000 enterprise E2E tests, UI authentication is a major bottleneck. How do you design an authentication architecture in Playwright that minimizes pipeline runtime while keeping tests isolated?**

**Spoken Answer:**

"Running UI authentication across 5,000 tests is an anti-pattern. If each login takes 4 seconds, you waste over 5.5 hours of machine time just rendering login forms.

To resolve this, I implement a two-tiered session reuse strategy:

First, in `playwright.config.ts`, I define a global `setup` project that runs before any test projects. In that setup project, we authenticate **once**. Ideally, we bypass the UI entirely by calling the authentication REST API via `playwright.request.newContext()`, obtaining the JWT access and refresh tokens, and writing a `.auth/user.json` storage state file to disk. If multi-factor authentication or complex redirects require UI login, we perform that UI login once in the setup project and call `await page.context().storageState({ path: '.auth/user.json' })`.

Second, in our dependent browser projects, we inject that pre-authenticated storage state via the `storageState` property in the project configuration. When workers spawn, each `BrowserContext` is hydrated with these cookies and localStorage values in under 20 milliseconds, without ever touching the `/login` route.

For scenarios requiring different user roles—like Admin, Standard, and Read-Only—we matrix the setup project into `admin.setup.ts` and `user.setup.ts`, outputting separate storage state files that map cleanly to corresponding test projects. This turns a multi-hour authentication tax into a single 2-second setup step."

### Question 3: Explain Playwright's Strict Mode. Why is it enabled by default, and how should an SDET resolve strict mode violations cleanly?

#### Spoken Answer:

"Strict Mode enforces that any locator passed to an action or web-first assertion must resolve to **exactly one** DOM element. If a locator matches two or more elements, Playwright refuses to pick one arbitrarily; instead, it immediately halts and throws a `StrictnessError`, listing all matched elements and their generated selector suggestions.

This is enabled by default because in traditional tools like Selenium, `driver.findElement()` silently returns the first matching element in document order. If a page refactor adds a hidden mobile menu or a duplicate button earlier in the DOM, Selenium tests continue clicking the wrong element, resulting in silent failures or confusing downstream errors.

To resolve strict mode violations cleanly, I follow three rules:

1. **Do not use `.first()` or `.nth(0)` as a quick fix:** That merely masks locator ambiguity and reintroduces flakiness if DOM order changes.
2. **Scope by accessibility roles and labels:** Narrow the locator using user-facing constraints, such as `page.getByRole('button', { name: 'Save Invoice' })`.
3. **Chain with `locator.filter()`:** Filter the collection by structural context—for instance, locating a row by text: `page.getByRole('row').filter({ hasText: 'INV-1092' }).getByRole('button', { name: 'Delete' })`.

This ensures locators remain unambiguous, resilient to layout changes, and aligned with user accessibility semantics."

### Question 4: How does Playwright handle iframes, and why is `frameLocator` architecturally superior to Selenium's `driver.switchTo().frame()`?

#### Spoken Answer:

"In Selenium, iframe automation is imperative and stateful. You invoke `driver.switchTo().frame('payment-frame')`, which mutates the driver session's internal pointer. If the iframe re-renders due to a React/Vue state update, the driver holds a stale reference, throwing `StaleElementReferenceException`. Furthermore, you have to remember to switch back to the main document with `driver.switchTo().defaultContent()`.

Playwright solves this with **declarative, stateless `frameLocator` objects**:

1. **Stateless Scope:** `const frame = page.frameLocator('iframe#stripe');` does not mutate page state. You can interact with the frame and the main page interchangeably without switching contexts.
2. **Auto-Waiting Across Boundaries:** When you call `frame.getByRole('textbox', { name: 'Card Number' }).fill(...)`, Playwright automatically waits for the iframe element to enter the DOM, verifies the frame document has loaded, locates the inner element, applies all 6 actionability checks, and performs the fill.
3. **Resilience to Re-renders:** If the iframe unmounts and remounts dynamically, `frameLocator` re-resolves the frame locator on demand rather than failing with stale references.
4. **Clean Nesting:** Nested iframes can be chained cleanly:
 

```
page.frameLocator('#parent').frameLocator('#child').getByRole('button')."
```

**Question 5: What are the 6 actionability checks performed during `locator.click()`, and why is using `{ force: true }` generally considered an anti-pattern?**

**Spoken Answer:**

"Before executing a `click()`, Playwright's engine runs an automated actionability pipeline:

1. **Attached:** The target element is connected to the live DOM tree.
2. **Visible:** The element has a non-empty bounding box, is not `display: none`, and is not `visibility: hidden`.
3. **Stable:** The element's CSS transitions or JavaScript animations have completed. Playwright checks that the bounding box remains identical across two consecutive animation frames (`requestAnimationFrame`).
4. **Enabled:** The element does not have the `disabled` attribute and is not within a disabled `<fieldset>`.
5. **Editable:** Applicable to text fields—ensures the element is not `readonly`.
6. **Receives Events (Hit-Testing):** Playwright computes the element's center point (`x, y`) and calls `document.elementFromPoint(x, y)` to confirm that the hit-test target resolves directly to the element, and is not obscured by a floating modal, tooltip, or loading overlay.

Using `{ force: true }` bypasses checks 2 through 6. Instead of simulating real user input, it dispatches a synthetic JavaScript click event directly to the node. This is a dangerous anti-pattern because if a loading spinner or modal backdrop covers the button, a real user cannot click it. Forcing the click produces a false positive, allowing tests to pass in CI while real users encounter a blocked interface in production."

**Question 6: What is the architectural difference between synchronous assertions (e.g., `expect(await locator.isVisible()).toBe(true)`) and Web-First Assertions (`await expect(locator).toBeVisible()`)?**

**Spoken Answer:**

"The difference comes down to **point-in-time snapshot evaluation versus continuous polling**.

When you write `expect(await locator.isVisible()).toBe(true)`:

1. `await locator.isVisible()` executes once immediately.
2. If the frontend framework (React, Angular) is in the middle of a microtask cycle and takes 50 milliseconds to mount that element, `isVisible()` returns `false`.
3. The assertion resolves to Jest's synchronous `expect(false).toBe(true)`, which fails instantly. Engineers often misdiagnose this as an infrastructure issue and add arbitrary sleeps.

In contrast, `await expect(locator).toBeVisible()` is a **Web-First Assertion**:

1. It initiates an internal polling loop inside the browser context that checks the element continuously.
2. If the element is absent, hidden, or animating, Playwright backs off, waits a few milliseconds, and re-evaluates.
3. It continues polling until the condition is met or the `expect.timeout` (default 5 seconds) expires.

Web-First Assertions eliminate race conditions, handle asynchronous rendering cleanly, and automatically handle inverted conditions—like waiting for spinners to disappear using `await expect(spinner).not.toBeVisible()`.

## Question 7: How do you implement network interception to mock third-party dependencies, test error scenarios, and speed up CI execution?

### Spoken Answer:

"We use `page.route()` or `context.route()` to intercept HTTP and HTTPS traffic directly at the browser network layer before requests go over the wire.

I use network routing for three primary enterprise use cases:

1. **Chaos Testing and Fault Tolerance:** We can simulate backend outages and verify UI error boundaries by fulfilling requests with error status codes:

```
await page.route('**/api/v1/payments', route => {
  route.fulfill({
    status: 503,
    contentType: 'application/json',
    body: JSON.stringify({ error: 'Gateway Timeout' })
  });
});
```

2. **Payload Mutation:** Instead of building complete mock backends, we can intercept real API responses, modify specific fields, and return the mutated data to the browser:

```
await page.route('**/api/v1/user', async route => {
  const response = await route.fetch();
  const json = await response.json();
  json.roles = ['SUPER_ADMIN'];
  await route.fulfill({ response, json });
});
```

3. **Third-Party Telemetry Blocking:** Third-party trackers like Google Analytics, Datadog, and Hotjar add significant network overhead and can delay `load` events. We abort them using regex filters:

```
await page.route(/(google-analytics|hotjar|datadog)/, route => route.abort());
```

This alone often reduces CI execution times by 30 to 40%."

## Question 8: How do Playwright fixtures work under the hood, and how are they superior to TestNG/JUnit `@BeforeMethod` and `@AfterMethod` annotations?

### Spoken Answer:

"Playwright fixtures use a **hierarchical Dependency Injection (DI) system** driven by generator-like execution contexts ( `test.extend()` ).

In TestNG or JUnit, lifecycle methods run imperatively. Test classes rely on mutable shared state—such as `private WebDriver driver;`—initialized in `@BeforeMethod` and torn down in `@AfterMethod` . If an engineer forgets a try-catch block, teardown logic fails, leaving zombie browser processes and leaking state into subsequent tests.

Playwright fixtures solve this with three architectural guarantees:

1. **Lazy Evaluation & Explicit Dependencies:** A fixture executes only if a test explicitly includes it in its argument list. If a test doesn't request `authenticatedPage` , that fixture never initializes.
2. **Enclosed Lifecycle Boundaries via `use()`** : Fixtures wrap test execution. Setup occurs before `await use(instance)` , and teardown logic runs immediately after `use()` completes. Even if a test times out or throws an unhandled exception, Playwright's runtime guarantees the post- `use()` teardown runs cleanly.
3. **Configurable Scoping:** Fixtures support two scopes: `test` scope (ephemeral, runs per test) and `worker` scope (runs once per OS worker process). This allows us to share heavy resources—like database connection pools or API clients—across tests in a worker while maintaining pristine test-level browser contexts."

## Question 9: Describe your strategy for debugging a flaky test that passes locally on macOS but intermittently fails in headless Linux CI environments.

### Spoken Answer:

"Flakiness that only manifests in headless Linux CI usually stems from one of four root causes: viewport differences, timing and CPU starvation, font rendering or subpixel shifts, or unhandled asynchronous race conditions.

My debugging playbook follows a structured triage process:

**1. Inspect the Playwright Trace ( `trace.zip` ):**

In our CI configuration, we set `trace: 'on-first-retry'` . When the test fails, Playwright captures a complete execution trace containing DOM snapshots before and after every action, network HAR logs, console logs, and actionability checks. I download the artifact and run `npx playwright show-trace trace.zip` .

**2. Inspect the Actionability Check Log:**

Trace Viewer reveals the exact check that stalled. For example, if `click()` timed out, the log might show: `waiting for element to receive pointer events – <div class="toast-notification"> intercepts pointer events` . This immediately tells me a notification overlay was covering the element on slower CI runners.

**3. Check Network Activity:**

The embedded network tab shows whether an asynchronous API call was pending or returned a `429 Too Many Requests` due to parallel workers overwhelming a shared staging gateway.

**4. Reproduce Locally with Linux Docker Containers:**

If local reproduction fails, I run the test inside Playwright's official Docker container (`mcr.microsoft.com/playwright` ) with `--headed` or `--debug` flags enabled, and throttle CPU to simulate resource-constrained CI runners:

```
docker run -it --rm --ipc=host -v $(pwd):/work/ -w /work/ mcr.microsoft.com/playwright
```

**5. Fix the Root Cause:**

I replace any instantaneous assertions or layout-dependent selectors with accessible role-based locators and web-first assertions."

---

**Question 10: How do you design and execute horizontal CI sharding for a suite of 10,000 Playwright tests, and how do you aggregate the test results into a unified report?**

**Spoken Answer:**

"Running 10,000 end-to-end tests sequentially or on a single VM is impractical. We shard the test execution across multiple CI machines using Playwright's native sharding feature:

**1. Horizontal Test Sharding:**

In our CI workflow (such as GitHub Actions or GitLab CI), we define a matrix job with `N` runners (e.g., 10 parallel machines). Each runner executes:

```
npx playwright test --shard=${{ matrix.shardIndex }}/${{ matrix.shardTotal }} --repor
```

Playwright partitions the test files evenly across the runners based on file paths and execution history.

**2. Blob Reporting:**

Instead of having each shard generate a standalone HTML report, each runner writes a lightweight binary `blob` report into a shared output directory.

### 3. Consolidated Report Synthesis Job:

We configure a downstream aggregation job that runs after all shards complete ( `if: always()` ):

- It downloads the blob report artifacts from all 10 shards into a single directory.
- It runs Playwright's report merge tool:

```
npx playwright merge-reports --reporter=html ./all-blob-reports
```

- This merges all results, screenshots, videos, and retry histories into a single, unified HTML report published to GitHub Pages or AWS S3.

### 4. Balancing and Worker Optimization:

On each runner, we enable `fullyParallel: true` and configure workers based on available virtual CPUs ( `workers: '100%` ). We combine this with API data pre-seeding and storage-state caching to achieve linear scalability and keep total suite execution under 15 minutes."